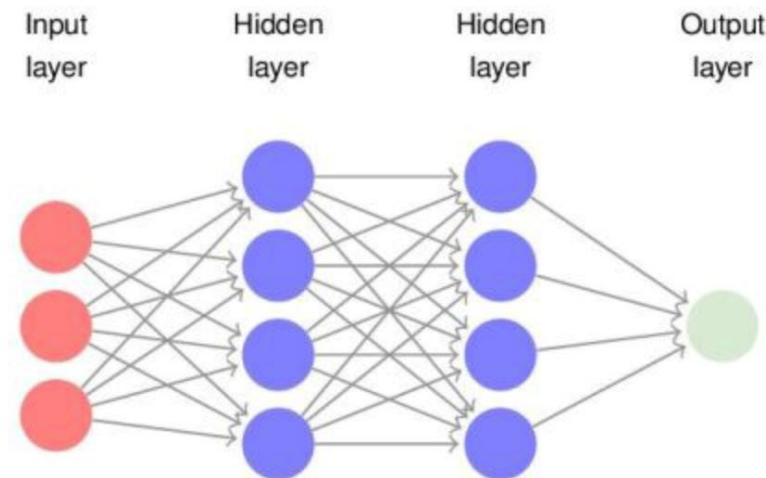

Deep Learning

Winter 2026

Neural Networks

Terminology



- Terms: Input layer, hidden layers, output layer
- A single hidden layer (of arbitrary width) can already be shown to be a **universal function approximator**
- Non-linear functions
 - are called activation functions in hidden layers
 - predict in the output layer and are used for the loss function

Feed Forward Neural Network: Basic Formulation

A feed forward neural network can be expressed by the following equation:

$$\hat{y}_k = g \left(\sum_{j=0}^M w_{jk} f \left(\sum_{i=0}^N w_{ij} x_i + b_j \right) + b_k \right)$$

Feed Forward Neural Network: Basic Formulation

$$\hat{y}_k = g \left(\sum_{j=0}^M w_{jk} f \left(\sum_{i=0}^N w_{ij} x_i + b_j \right) + b_k \right)$$

Diagram illustrating the basic formulation of a feed-forward neural network:

- real-value network output**: \hat{y}_k
- activation function**: g (e.g., sigmoid, tanh(), ReLU, etc.)
- input data**: x_i
- input-to-hidden weight matrix**: w_{ij} (i rows by j columns)
- input bias term**: b_j
- hidden-layer activation**: f
- number of inputs per sample**: N
- output bias term**: b_k
- hidden-layer activation**: f
- number of hidden layer units**: M
- number of outputs**: k

Feed Forward Neural Network: Basic Formulation

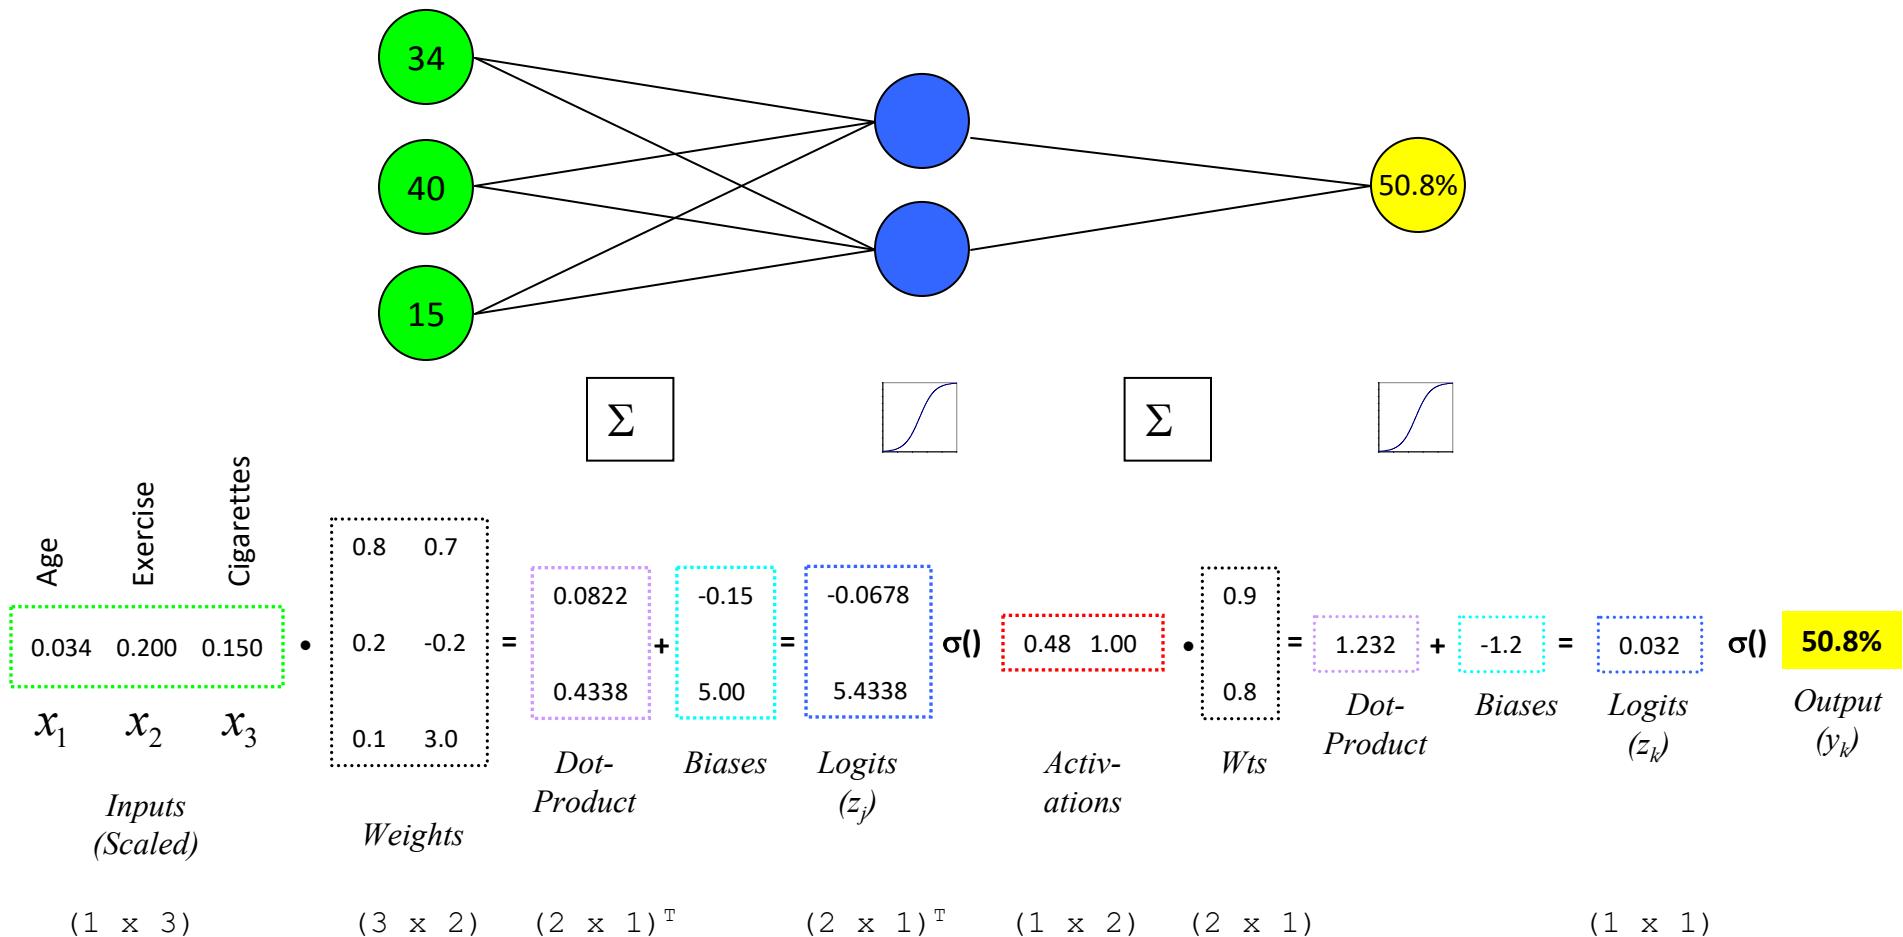
... or more compactly in Linear Algebra form:

$$\hat{\mathbf{y}} = \sigma(\mathbf{W}_2 \cdot \tanh(\mathbf{W}_1 \cdot \mathbf{x} + \mathbf{b}_1) + \mathbf{b}_2)$$

$$\hat{y}_k = g\left(\sum_{j=0}^M w_{jk} f\left(\sum_{i=0}^N w_{ij} x_i + b_j \right) + b_k \right)$$

An Example

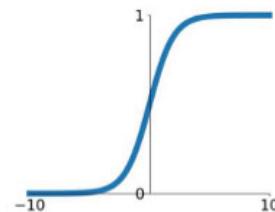
Consider a simple neural network architecture that computes the probability of being alive at the end of five years given “life style” features:



Activation Functions

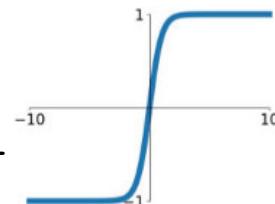
Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$



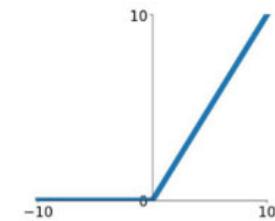
tanh

$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$



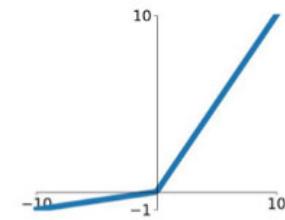
ReLU

$$\max(0, x)$$



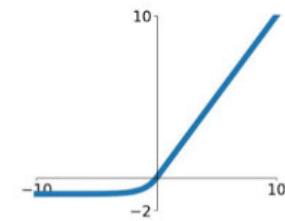
Leaky ReLU

$$\max(0.1x, x)$$



ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



Training

Step 1: The neural network is presented with a series input/output pairs $\langle \bar{x}_s, y_s \rangle \in S$ from the training set.

Step 2: The neural network is used to perform inference (a forward pass) to generate an estimated output \hat{y}_s .

Step 3: A cost function (see next slides) is used to quantify how well (or poorly) the neural network performed by comparing \hat{y}_s and y_s .

Step 4: Back propagation (see next slides) is performed by calculating gradients with respect to each trainable variable, which are used to update trainable variables:

$$\hat{y}_k = g\left(\sum_{j=0}^M w_{jk} f\left(\sum_{i=0}^N w_{ij} x_i + b_j \right) + b_k \right)$$

The updates can be performed as:

- Stochastic Gradient Descent: Network parameters are updated for each training example, or
- Batch Gradient Descent: Gradients are accumulated over the entire training set and applied all at once.

Cost Functions

Certain cost functions are associated with typical output layers, but they are unconstrained and are influenced by the neural network architecture:

- Regression uses Mean Square Error (MSE):

$$J = \frac{1}{2} \sum_s (\hat{y}_s - y_s)^2$$

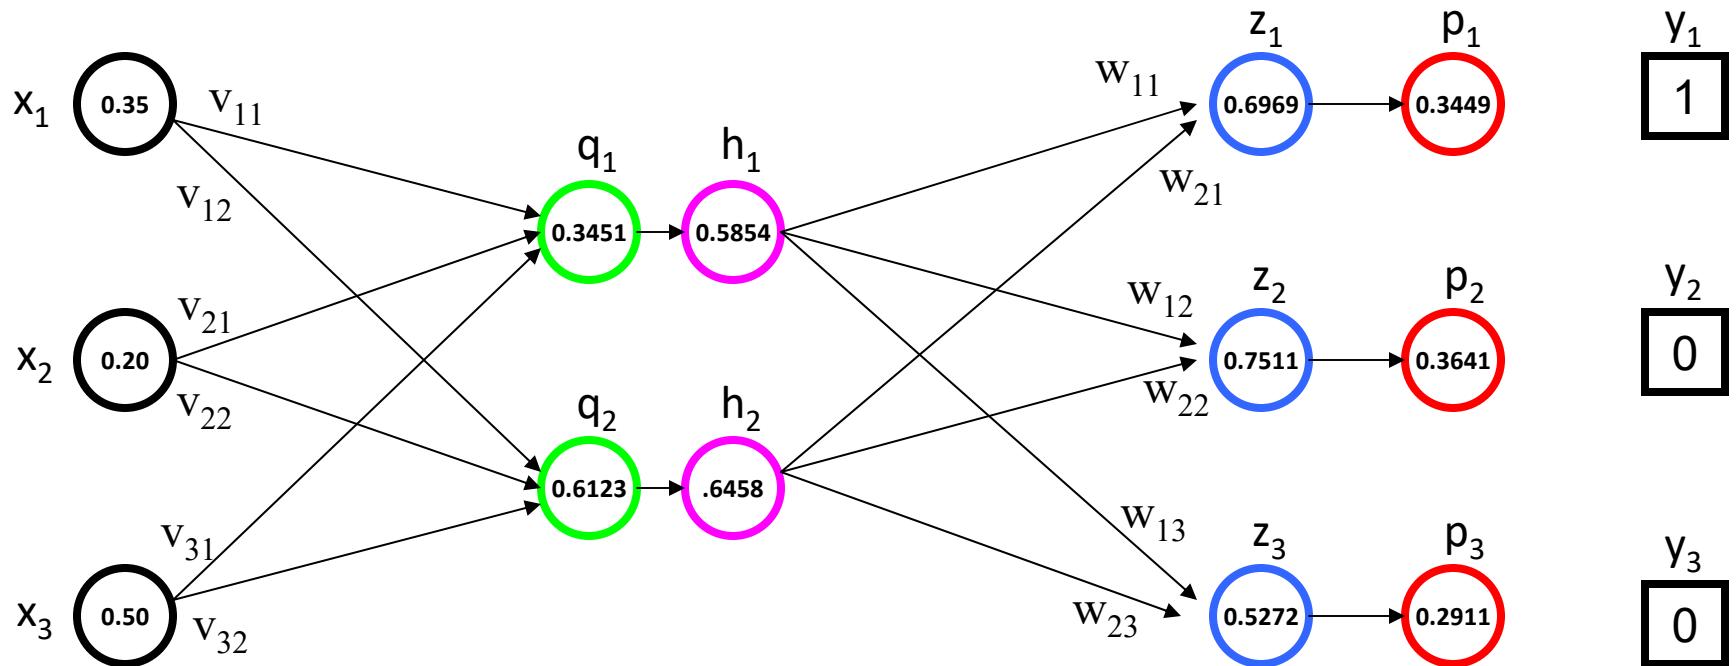
- Binary-Classification uses two-class cross-entropy:

$$J = -(y_s \ln(\hat{p}_s) + (1-y_s) \ln(1-\hat{p}_s))$$

- Multi-Class Classification uses multi-class cross-entropy:

$$J = -\sum_{c \in C} y_c^{(s)} \ln(\hat{p}_c^{(s)})$$

Gradient Descent -- Architecture



$$\sigma(q) = \frac{1}{1 + e^{-q}}$$

$$p_k = \frac{e^{z_k}}{\sum_{c \in C} e^{z_c}}$$

0.3500
0.2000
0.5000

0.4963	0.7682
0.0885	0.1320
0.3074	0.6341

0.3451
0.6123

0.5854
0.6485

0.4901	0.8964	0.4556
0.6323	0.3489	0.4017

0.6969
0.7511
0.5272

34.5%
36.4%
29.1%

Good
Neutral
Bad

Inputs X	Weights V	Linear Q	Hidden H	Weights W	Logits Z	Probs P	Labels Y
(1 \times 3)	(3 \times 2)	(1 \times 2)	(1 \times 2)	(2 \times 3)	(1 \times 3)	(1 \times 3)	(1 \times 3)

Gradient Descent -- Equations

Cost Function: Multi-Class Cross-Entropy

$$L = -\sum_{i=0}^C y_i \ln(p_i)$$

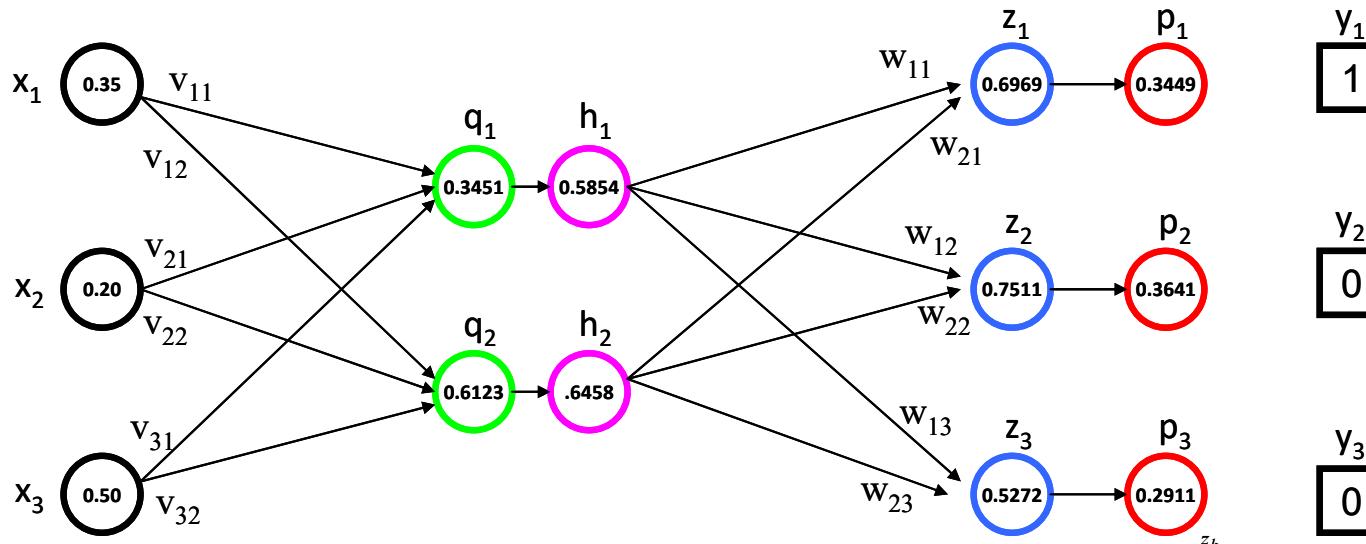
Activation Function: Sigmoid

$$\sigma(q) = \frac{1}{1 + e^{-q}}$$

Softmax:

$$\begin{aligned} p_k &= \frac{e^{z_k}}{\sum_{c \in C} e^{z_c}} \\ &= \frac{e^{z_1}}{e^{z_1} + e^{z_2} + e^{z_3}} \end{aligned}$$

Gradient Descent -- Forward Pass Calculations



$$\sigma(q) = \frac{1}{1 + e^{-q}}$$

$$p_k = \frac{e^{z_k}}{\sum_{c \in C} e^{z_c}}$$

0.3500
0.2000
0.5000

0.4963 0.7682
0.0885 0.1320
0.3074 0.6341

0.3451
0.6123

0.5854
0.6458

0.4901 0.8964 0.4556
0.6323 0.3489 0.4017

0.6969
0.7511
0.5272

34.5%
36.4%
29.1%

Good
Neutral
Bad

Inputs \$X\$
(1 x 3)

Weights \$V\$
(3 x 2)

Linear \$Q\$
(1 x 2)

Weights \$W\$
(2 x 3)

Logits \$Z\$
(1 x 3)

Probs \$P\$
(1 x 3)

Labels \$Y\$
(1 x 3)

$$\begin{bmatrix} 0.3500 & 0.2000 & 0.5000 \end{bmatrix} \bullet$$

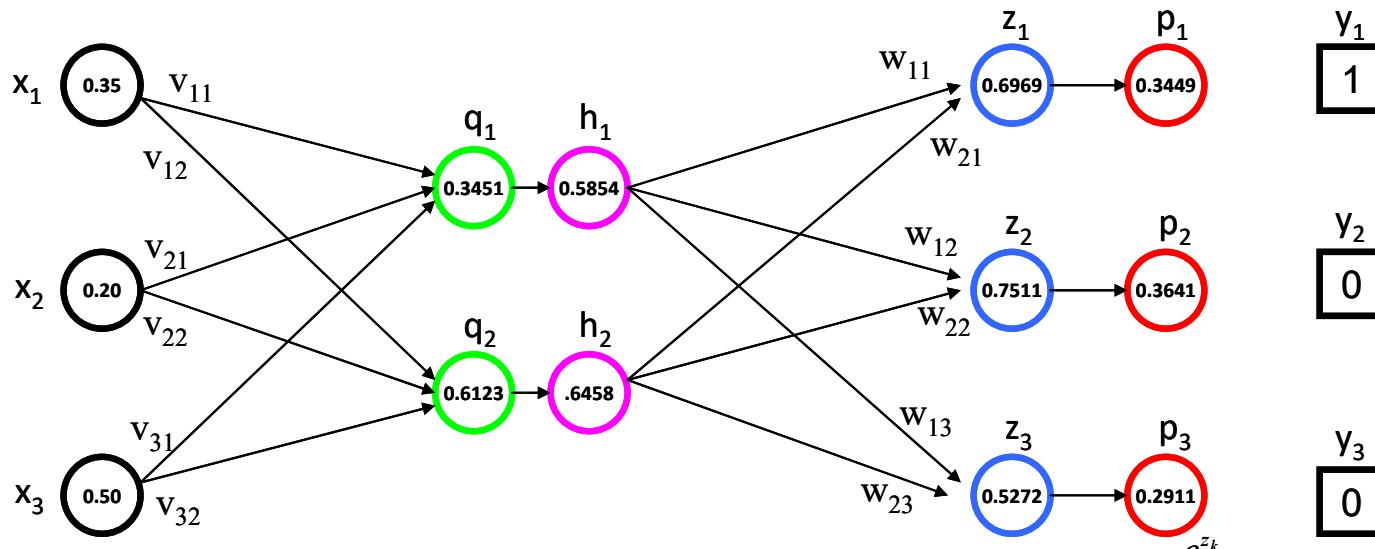
$$\begin{bmatrix} 0.4963 & 0.7682 \\ 0.0885 & 0.1320 \\ 0.3074 & 0.6341 \end{bmatrix}$$

=

$$\begin{bmatrix} 0.3451 \\ 0.6123 \end{bmatrix}$$

$$0.4963 \times 0.3500 + 0.0885 \times 0.2000 + 0.3074 \times 0.500 = 0.3451$$

Gradient Descent -- Forward Pass Calculations



$$\sigma(q) = \frac{1}{1 + e^{-q}}$$

$$p_k = \frac{e^{z_k}}{\sum_{c \in C} e^{z_c}}$$

0.3500
0.2000
0.5000

0.4963 0.7682
0.0885 0.1320
0.3074 0.6343

0.3451
0.6123

0.5854
0.6458

0.4901 0.8964 0.4556
0.6323 0.3489 0.4017

0.6969
0.7511
0.5272

34.5%
36.4%
29.1%

Good
Neutral
Bad

Inputs \$X\$
(1 x 3)

Weights \$V\$
(3 x 2)

Linear \$Q\$
(1 x 2)

Hidden \$H\$
(1 x 2)

Weights \$W\$
(2 x 3)

Logits \$Z\$
(1 x 3)

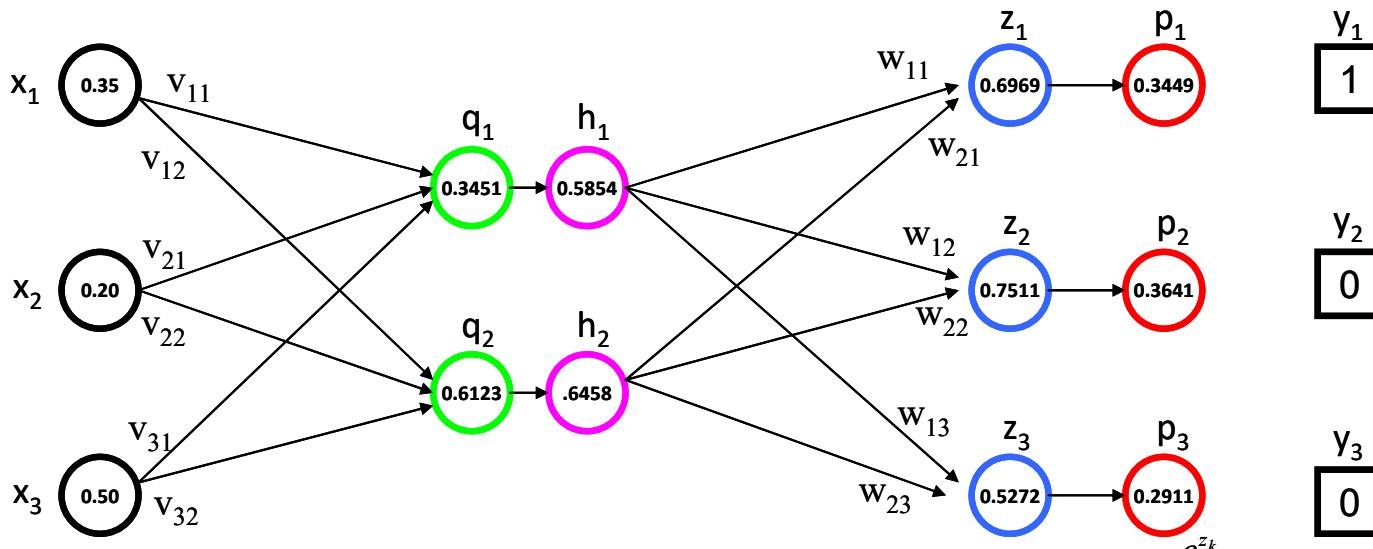
Probs \$P\$
(1 x 3)

Labels \$Y\$
(1 x 3)

$$\sigma \left(\begin{matrix} 0.3451 \\ 0.6123 \end{matrix} \right) = \begin{matrix} 0.5854 \\ 0.6458 \end{matrix}$$

$$1 / (1 + \exp(-0.3451)) = 0.5854$$

Gradient Descent -- Forward Pass Calculations



$$\sigma(q) = \frac{1}{1 + e^{-q}}$$

$$p_k = \frac{e^{z_k}}{\sum_{c \in C} e^{z_c}}$$

0.3500
0.2000
0.5000

0.4963	0.7682
0.0885	0.1320
0.3074	0.6341

0.3451
0.6123

0.58
0.64

0.4901	0.8964	0.4556
0.6323	0.3489	0.4017

0.6969
0.7511
0.5272

34.5%
36.4%
29.1%

Good
Neutral
Bad

Inputs X

Weights V
(3 x 2)

$$\begin{array}{cc} \textit{Linear } Q & \textit{Hidden } H \\ (1 \times 2) & (1 \times 2) \end{array}$$

Weights W
(2 x 3)

Logits
(1 x)

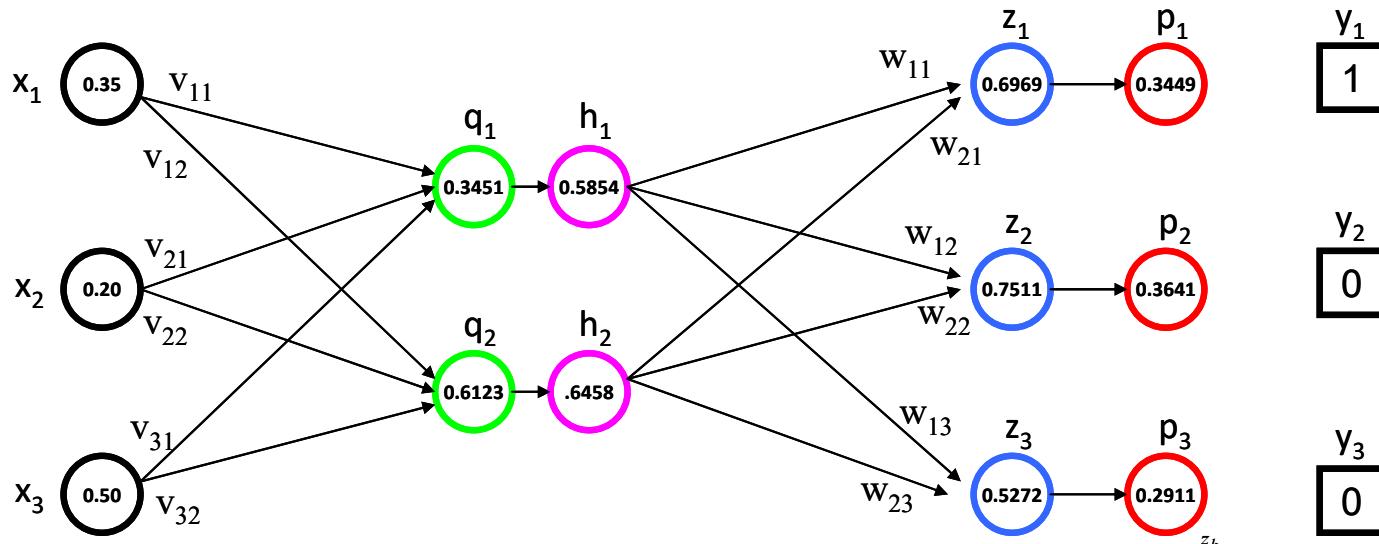
Probs P

Labels Y
(1 x 3)

$$\begin{array}{c} \boxed{0.5854 \quad 0.6485} \\ \bullet \quad \boxed{\begin{array}{ccc} 0.4901 & 0.8964 & 0.4556 \\ & \dots & \\ 0.6323 & 0.3489 & 0.4017 \end{array}} \\ = \quad \boxed{\begin{array}{c} 0.6969 \\ 0.7511 \\ 0.5272 \end{array}} \end{array}$$

$0.4901 \times 0.5854 + 0.6323 \times 0.6458 = 0.6969$

Gradient Descent -- Forward Pass Calculations



$$\sigma(q) = \frac{1}{1 + e^{-q}}$$

$$p_k = \frac{e^{z_k}}{\sum_{c \in C} e^{z_c}}$$

0.3500
0.2000
0.5000

0.4963	0.7682
0.0885	0.1320
0.3074	0.6341

0.3451
0.6123

0.5854
0.6458

0.4901	0.8964	0.4556
0.6323	0.3489	0.4017

0.6969
0.7511

34.5%
36.4%

Good
Neutral
Bad

Inputs \$X\$
(1 x 3)

Weights \$V\$
(3 x 2)

Linear \$Q\$
(1 x 2)

Hidden \$H\$
(1 x 2)

Weights \$W\$
(2 x 3)

Logits \$Z\$
(1 x 3)

Probs \$P\$
(1 x 3)

Labels \$Y\$
(1 x 3)

softmax(
0.6969
0.7511
0.5272
)

=

34.5%
36.4%
29.1%

$$\exp(0.6969)/[\exp(0.6969) + \exp(0.7511) + \exp(0.5272)] = 0.3449$$

Gradient Descent -- Some Derivatives

Cost Function: Multi-Class Cross-Entropy

$$\frac{\partial L}{\partial p_i} = -\frac{y_i}{p_i}$$

Activation Function: Sigmoid

$$\frac{\partial \sigma(q)}{\partial q} = \sigma(q)(1 - \sigma(q))$$

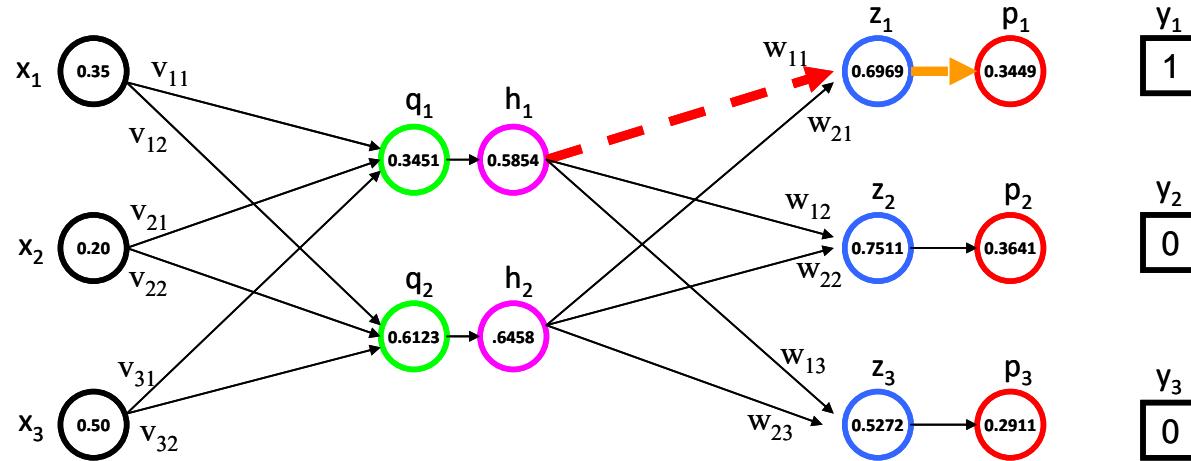
Softmax:

$$\frac{\partial \text{softmax}(z_k)}{\partial z_k} = \frac{e^{z_k} (e^{z_0} + e^{z_1} + \dots + e^{z_c} - e^{z_k})}{(e^{z_0} + e^{z_1} + \dots + e^{z_c})^2}$$

Linear Transform:

$$\frac{\partial w \cdot x}{\partial w} = x$$

Gradient Descent -- Gradient Calculation

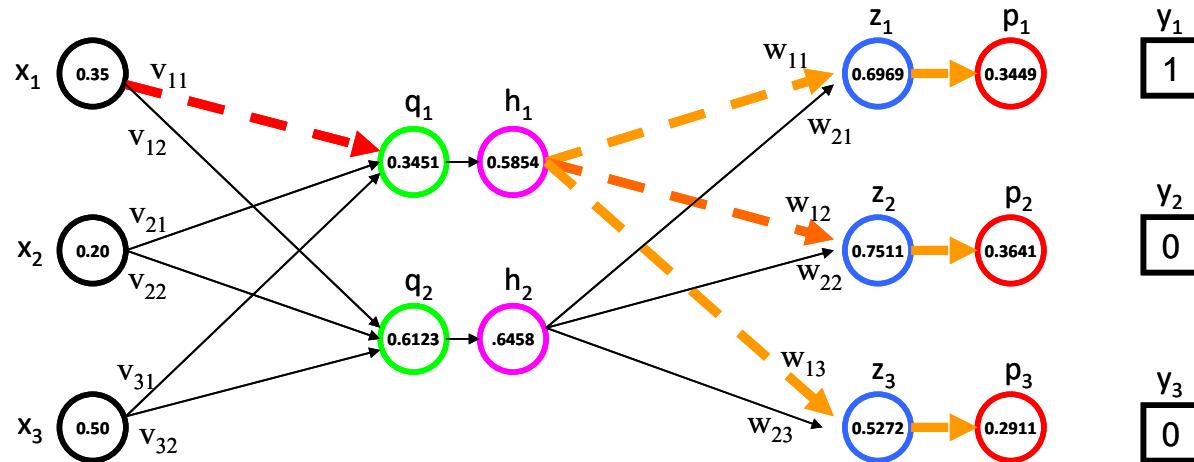


Chain Rule:

$$\frac{\partial f(g(x))}{\partial x} = \frac{\partial f(g(x))}{\partial g(x)} \cdot \frac{\partial g(x)}{\partial x}$$

$$\begin{aligned}
 \frac{\partial L}{\partial w_{11}} &= \frac{\partial L}{\partial p_1} \cdot \frac{\partial p_1}{\partial z_1} \cdot \frac{\partial z_1}{\partial w_{11}} \\
 &= -\frac{1}{p_1} \cdot \frac{e^{z_1} \cdot (e^{z_2} + e^{z_3})}{(e^{z_1} + e^{z_2} + e^{z_3})^2} \cdot h_1 \\
 &= -\frac{1}{0.3449} \cdot \frac{e^{0.6969} \cdot (e^{0.7511} + e^{0.5272})}{(e^{0.6969} + e^{0.7511} + e^{0.5272})^2} \cdot 0.5854 \\
 &= -0.3835
 \end{aligned}$$

Gradient Descent -- Gradient Calculation



$$\frac{\partial L}{\partial v_{11}} = \frac{\partial L}{\partial p_1} \cdot \frac{\partial p_1}{\partial z_1} \cdot \frac{\partial z_1}{\partial h_1} \cdot \frac{\partial h_1}{\partial q_1} \cdot \frac{\partial q_1}{\partial v_{11}} + \frac{\partial L}{\partial p_1} \cdot \frac{\partial p_1}{\partial z_2} \cdot \frac{\partial z_2}{\partial h_1} \cdot \frac{\partial h_1}{\partial q_1} \cdot \frac{\partial q_1}{\partial v_{11}} + \frac{\partial L}{\partial p_1} \cdot \frac{\partial p_1}{\partial z_3} \cdot \frac{\partial z_3}{\partial h_1} \cdot \frac{\partial h_1}{\partial q_1} \cdot \frac{\partial q_1}{\partial v_{11}}$$

Gradient Descent -- Apply Gradients

Initialize $\theta^{(0)}$

Repeat until convergence:

$$\theta^{(t+1)} = \theta^{(t)} - \eta \nabla_{\theta^{(t)}} L(X, Y; \theta^{(t)})$$

Return $\theta^{(t_{max})}$

$$\begin{array}{c} \begin{matrix} 0.4901 & 0.8964 & 0.4556 \\ 0.6323 & 0.3489 & 0.4017 \end{matrix} \\ \text{Weights } W^{(t)} \end{array} \quad - 0.5 \quad \begin{array}{c} \begin{matrix} -0.3835 & - & - \\ - & - & - \end{matrix} \\ \nabla_{w_{11}} L(X, Y, \theta^{(t)}) \end{array} = \begin{array}{c} \begin{matrix} 0.6819 & 0.8964 & 0.4556 \\ 0.6323 & 0.3489 & 0.4017 \end{matrix} \\ \text{Weights } W^{(t+1)} \end{array}$$

$$\begin{array}{c} \begin{matrix} 0.4963 & 0.7682 \\ 0.0885 & 0.1320 \\ 0.3074 & 0.6341 \end{matrix} \\ \text{Weights } V^{(t)} \end{array} \quad - 0.5 \quad \begin{array}{c} \begin{matrix} - & - \\ - & - \end{matrix} \\ \nabla_{v_{11}} L(X, Y, \theta^{(t)}) \end{array} = \begin{array}{c} \begin{matrix} 0.7682 \\ 0.0885 & 0.1320 \\ 0.3074 & 0.6341 \end{matrix} \\ \text{Weights } V^{(t+1)} \end{array}$$

Gradient Descent -- Additional Resources

Detailed Numerical Example:

<https://medium.com/@14prakash/back-propagation-is-very-simple-who-made-it-complicated-97b794c97e5c>

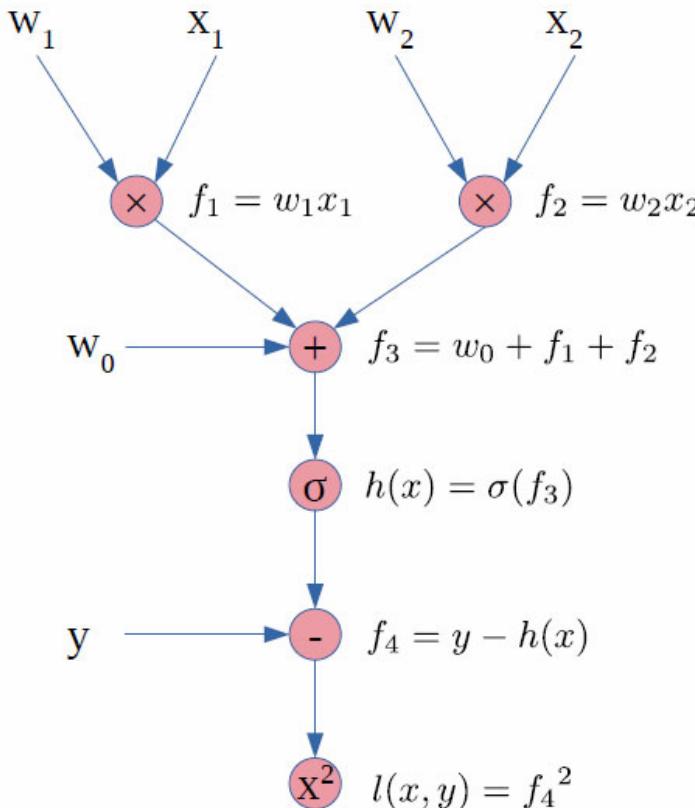
Gradient Calculation Video:

<https://www.youtube.com/watch?v=lj7JYJSgpZI>

Wolfram Alpha:

<https://www.wolframalpha.com/>

Calculation Graphs Forward Pass

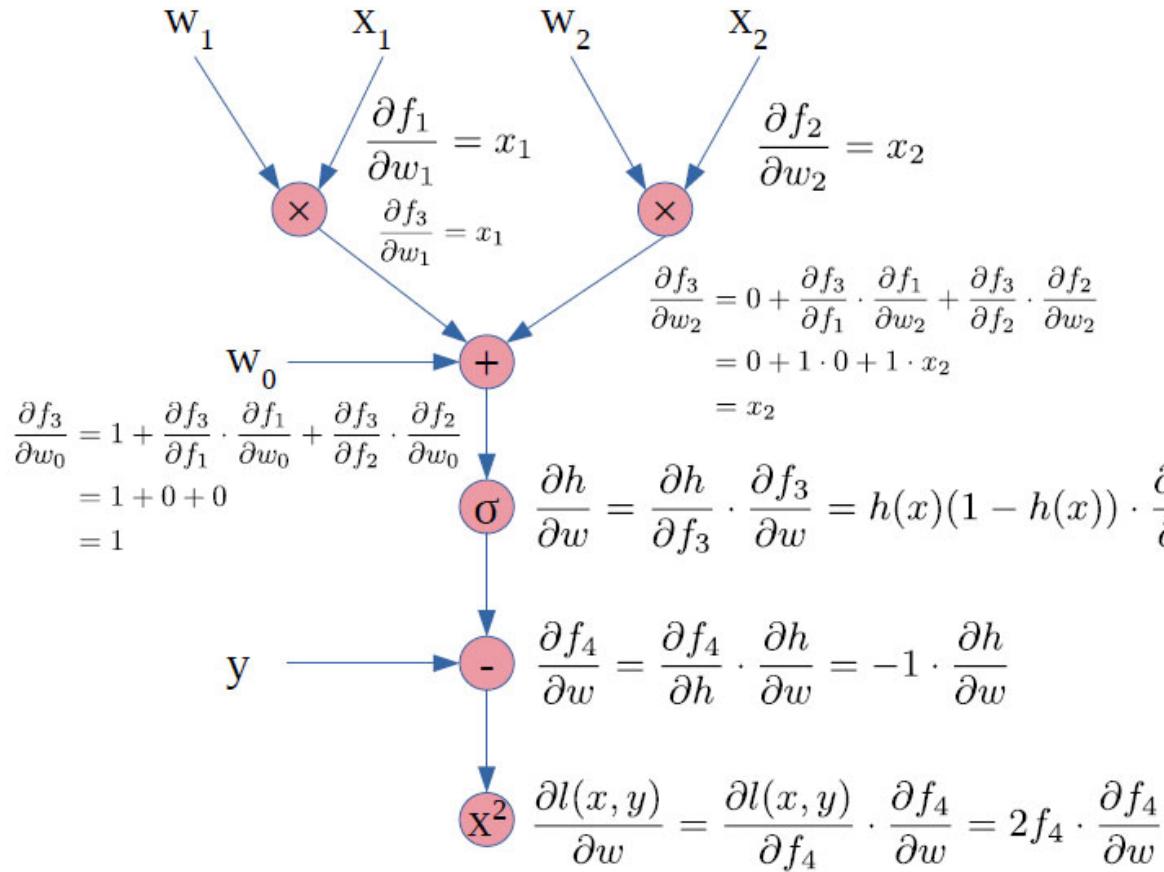


$$h(x) = \sigma(w_0 + w_1x_1 + w_2x_2)$$

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

$$l(x, y) = (y - h(x))^2$$

Calculation Graphs Gradients



Tunable Hyper-parameters

- Learning Rate is a “knob” that you can twist empirically
- Stochastic Gradient Descent vs. Batch Gradient Descent
- Drop Out randomly ignores certain weight dimensions during training
- Regularization minimizes a specific property of the neural network by modifying the cost function:

$$J = \frac{1}{2} \sum_S (\hat{y}_s - y_s)^2 + \lambda_1 \|w_{ij}\|_2 + \lambda_2 \|w_{jk}\|_2$$

- Scaling input values and weight initializations
- Gradient clipping and normalization
- Dynamic learning rate algorithms
- Many others