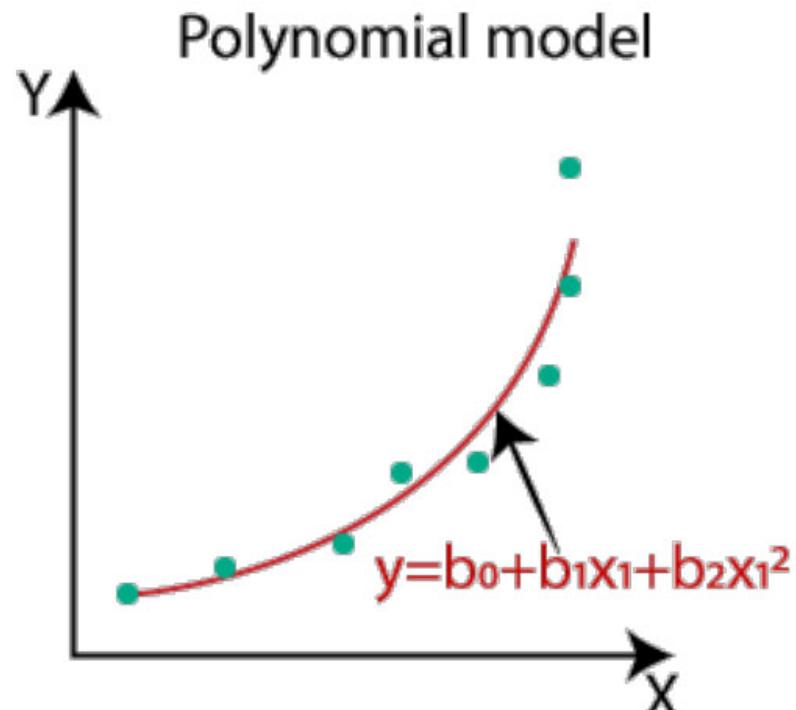
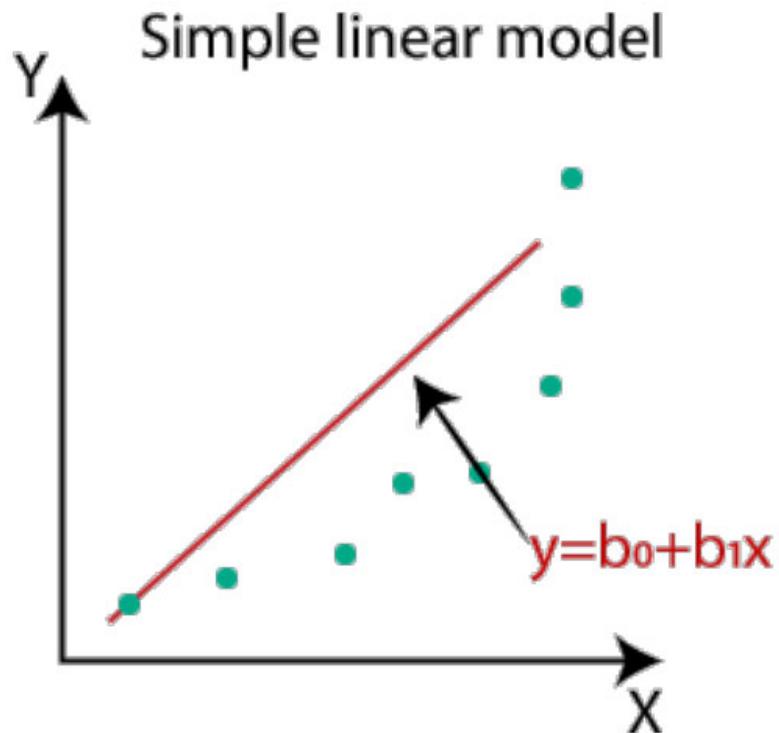

Deep Learning

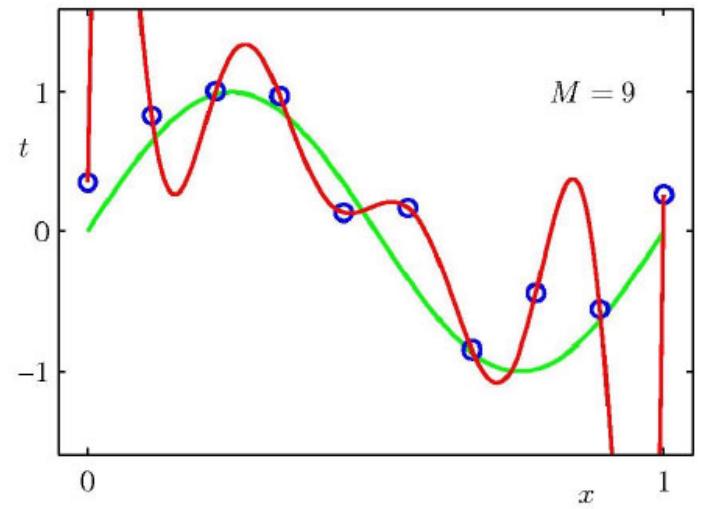
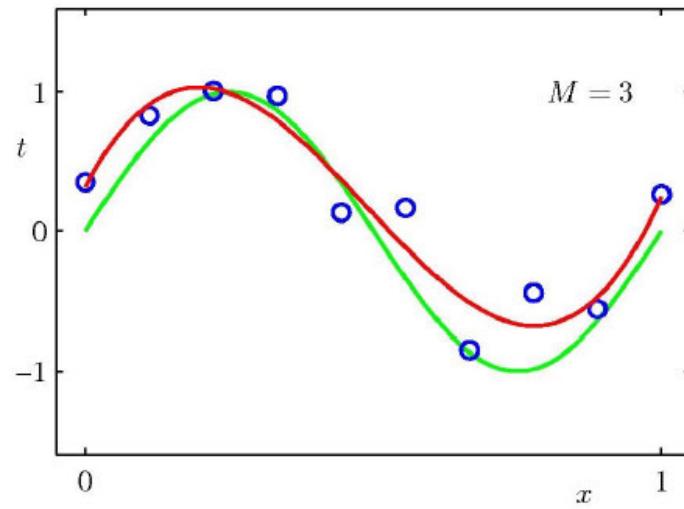
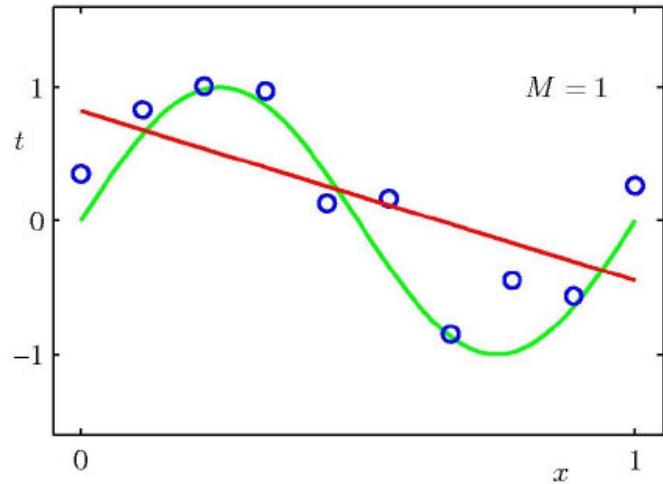
Winter 2026

Regularization

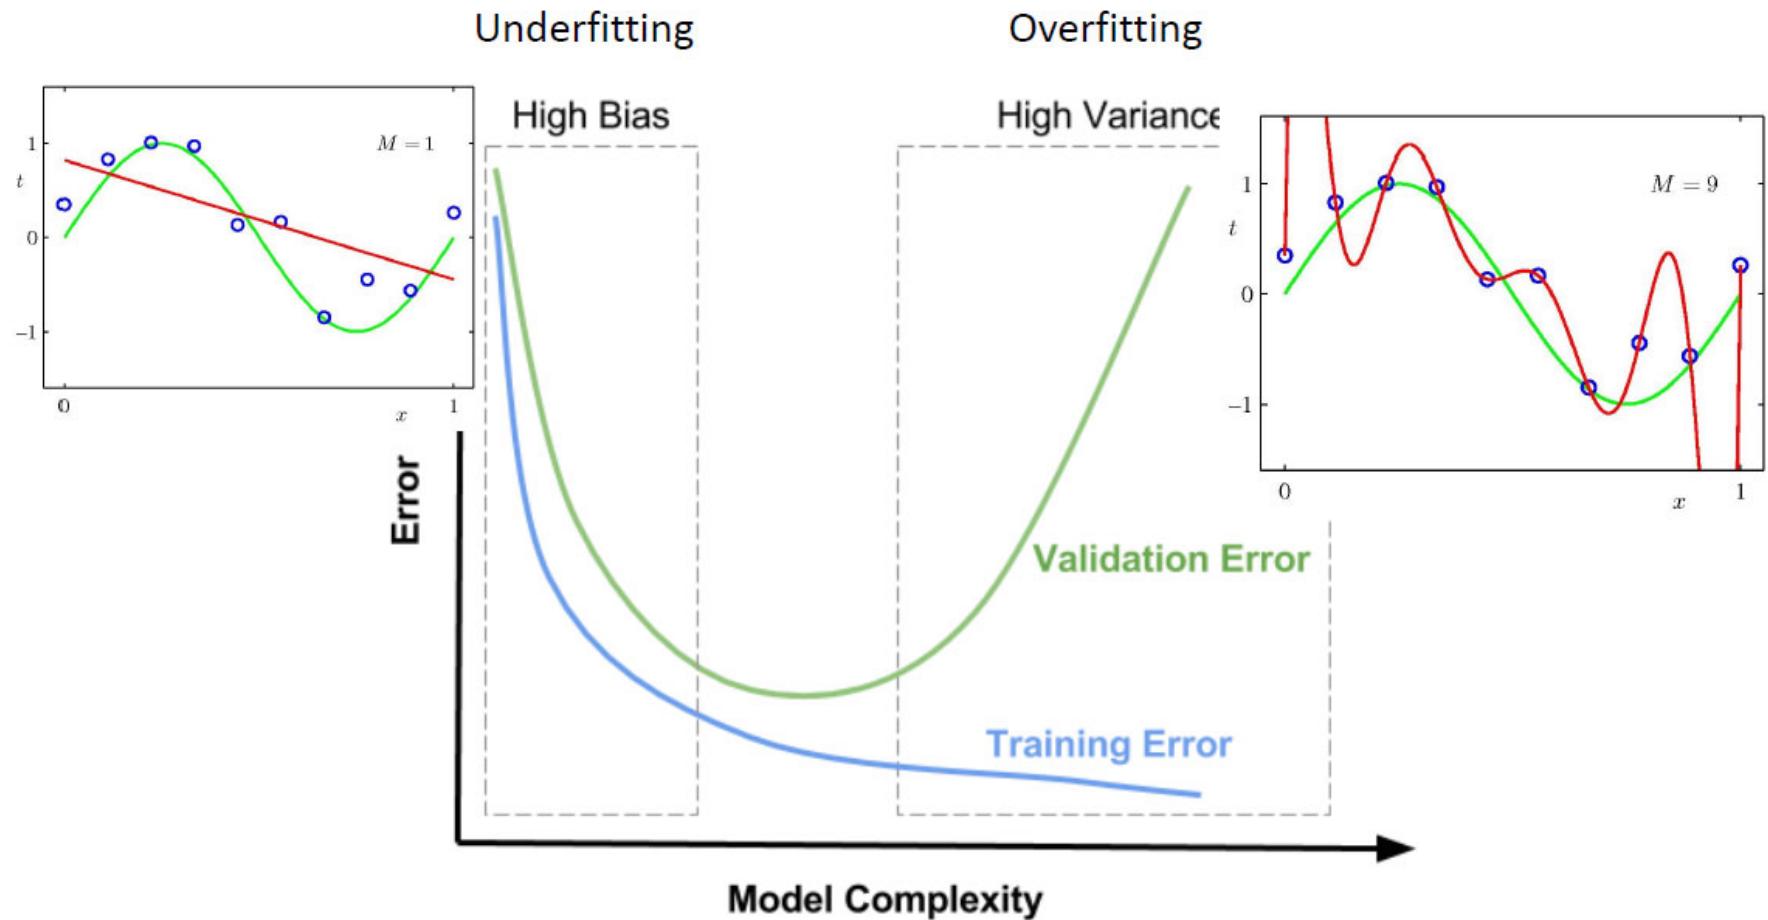
Polynomial Regression



Under/Over Fitting



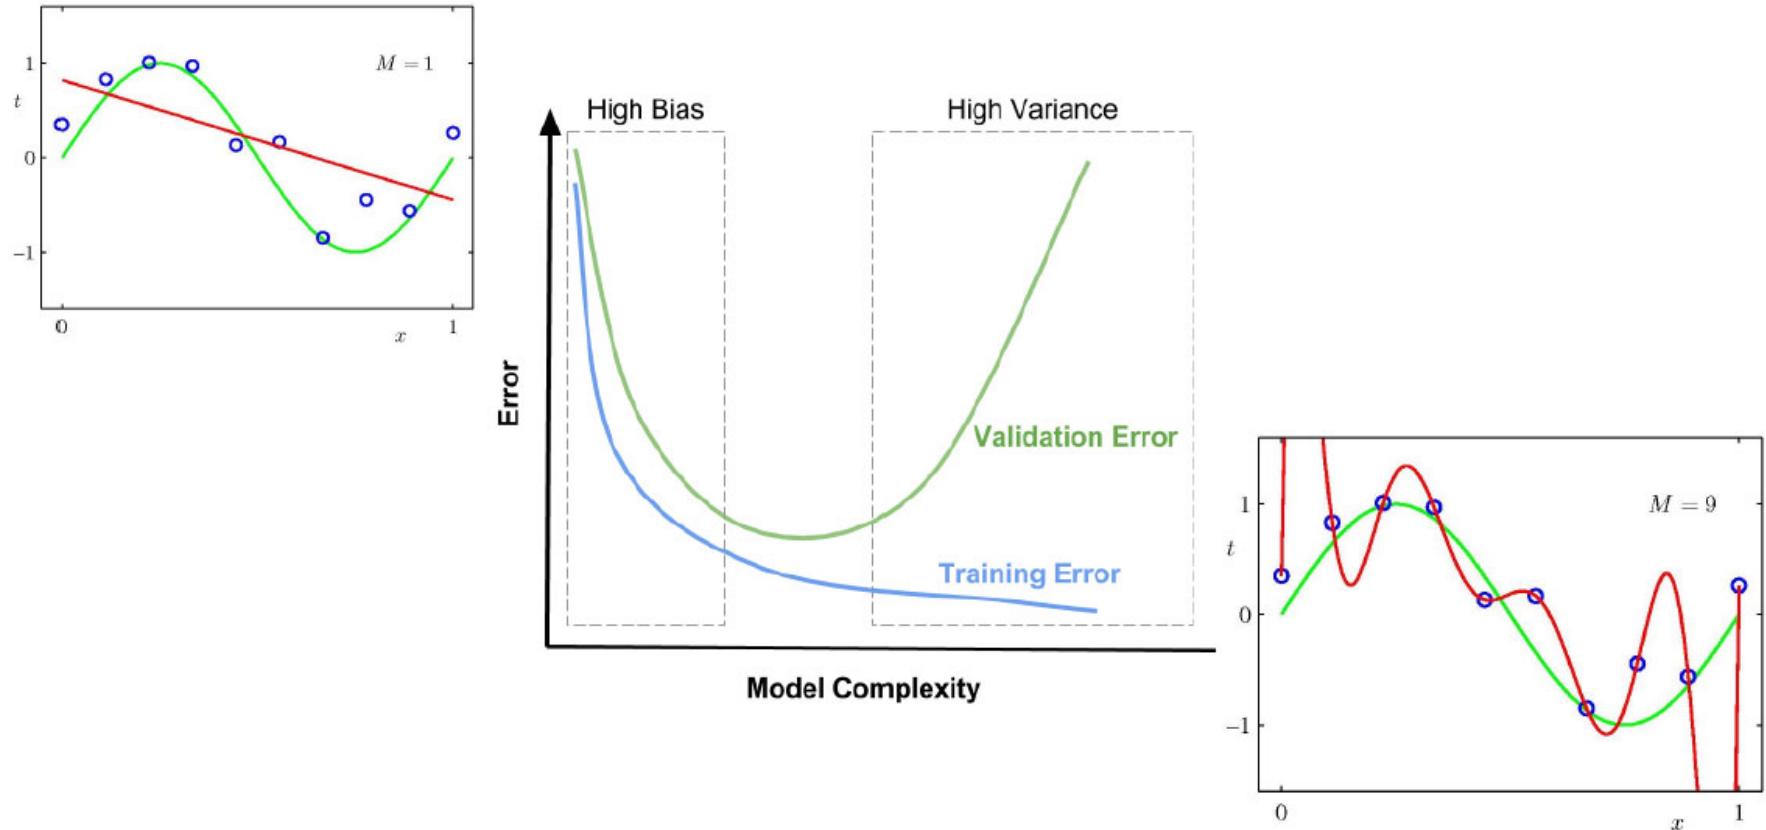
Bias/Variance Trade-off



What Goes Wrong?

- Underfit: too many assumptions
 - Hypothesis class doesn't contain the optimal hypothesis
 - (or even a “good” hypothesis)
 - Data representation discards essential information
- Overfit: not enough assumptions
 - Hypothesis class is “hard to search”
 - Learning algorithm is inefficient, can't optimize parameters
 - Many hypotheses perfectly fit training data

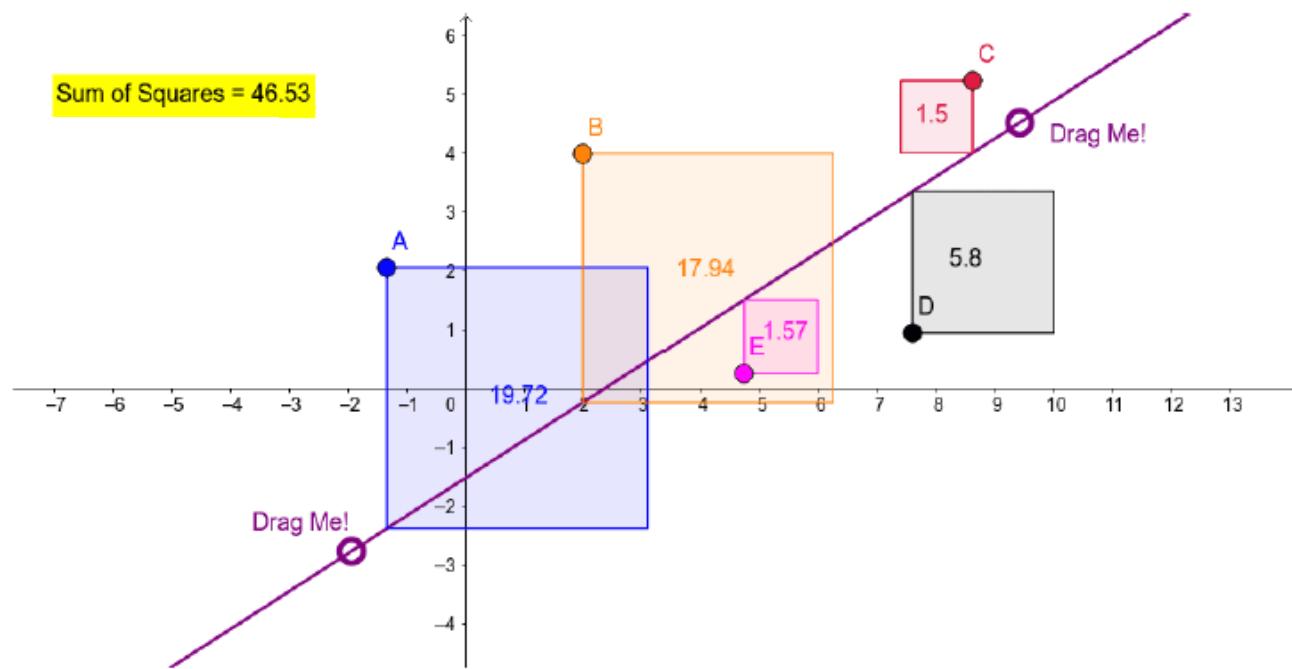
How Do We Stop Overfitting



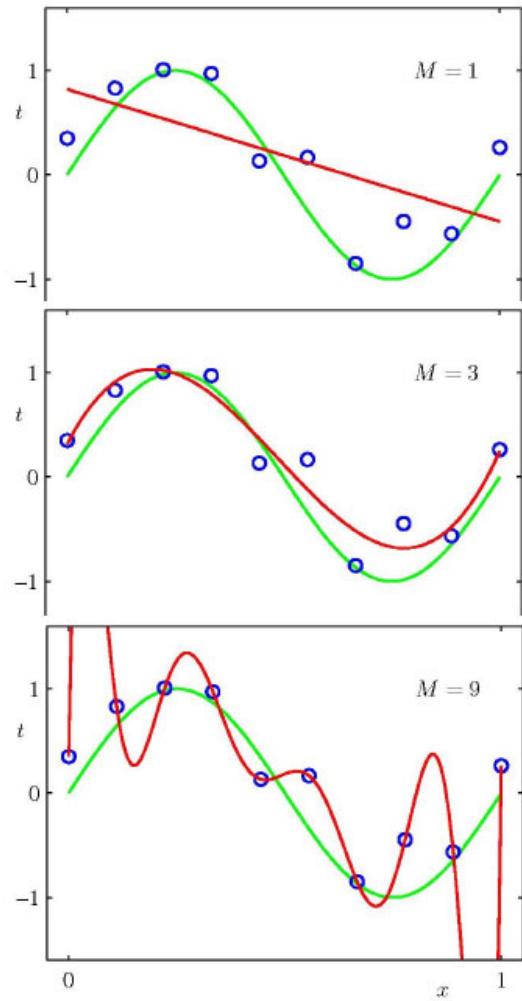
Motivating Regularization

- Our predictor f with parameters w has loss L :

$$L(y, f(x; w)) = (y - w \cdot x)^2$$



Motivating Regularization

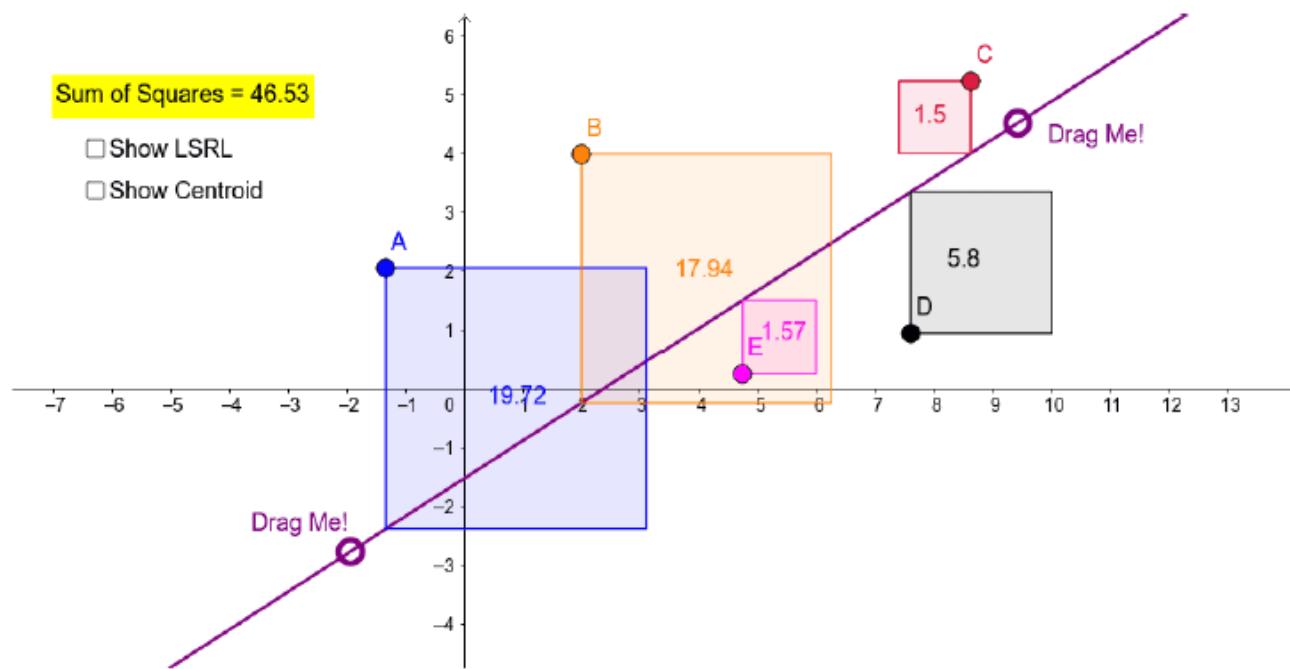


	$M = 1$	$M = 3$	$M = 9$
w_0^*	0.82	0.31	0.35
w_1^*	-1.27	7.99	232.37
w_2^*		-25.43	-5321.83
w_3^*		17.37	48568.31
w_4^*			-231639.30
w_5^*			640042.26
w_6^*			-1061800.52
w_7^*			1042400.18
w_8^*			-557682.99
w_9^*			125201.43

Motivating Regularization

- What if we add a term to our loss?

$$L(y, f(x; w)) = (y - w \cdot x)^2 + \lambda w^2$$



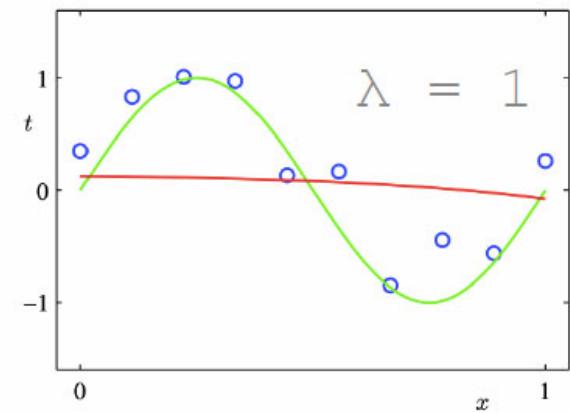
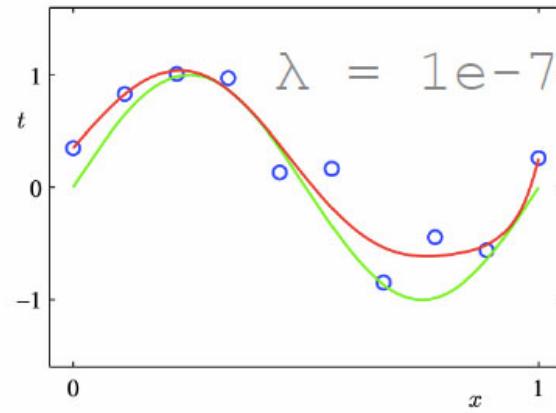
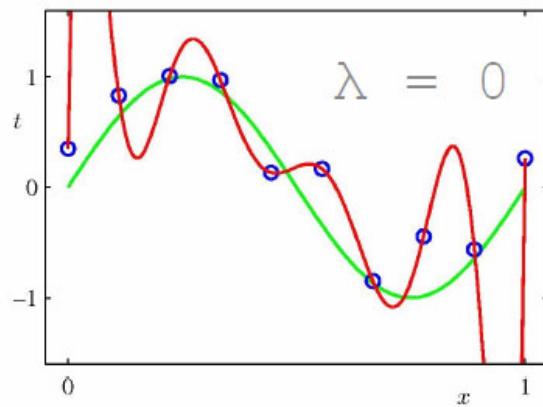
How Do We Stop Overfitting?

$$L(y, f(x; w)) = (y - w \cdot x)^2 + \lambda w^2$$

	$\lambda = 0$	$1e-7$	1
w_0^*	0.35	0.35	0.13
w_1^*	232.37	4.74	-0.05
w_2^*	-5321.83	-0.77	-0.06
w_3^*	48568.31	-31.97	-0.05
w_4^*	-231639.30	-3.89	-0.03
w_5^*	640042.26	55.28	-0.02
w_6^*	-1061800.52	41.32	-0.01
w_7^*	1042400.18	-45.95	-0.00
w_8^*	-557682.99	-91.53	0.00
w_9^*	125201.43	72.68	0.01

How Do We Stop Overfitting?

$$L(y, f(x; w)) = (y - w \cdot x)^2 + \lambda w^2$$



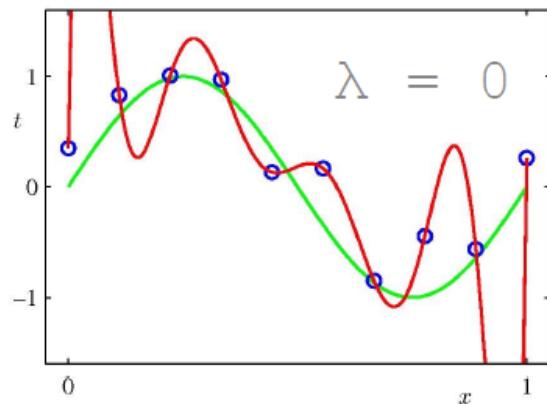
Regularized Least Squares

$$\mathcal{L}(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^N \{y_i - \mathbf{w}^T \cdot \mathbf{x}_i\}^2 \quad \mathcal{L}(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^N \{y_i - \mathbf{w}^T \cdot \mathbf{x}_i\}^2 + \lambda \frac{1}{2} \mathbf{w}^T \mathbf{w}$$

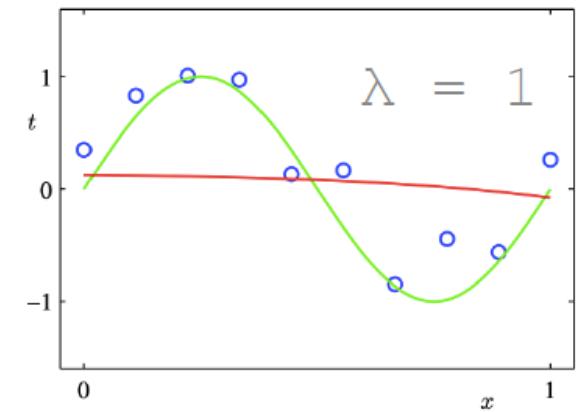
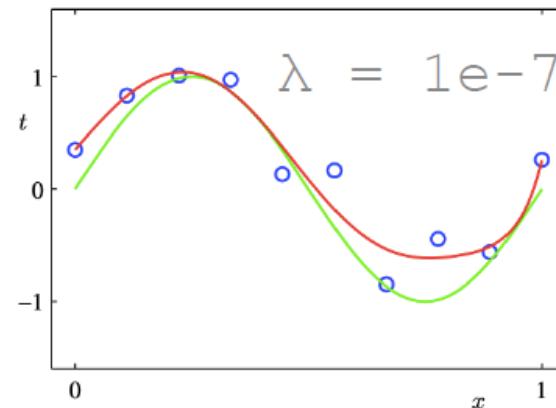
$$\mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y} \quad \mathbf{w} = (\lambda \mathbf{I} + \mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y}$$

How Do We Stop Overfitting?

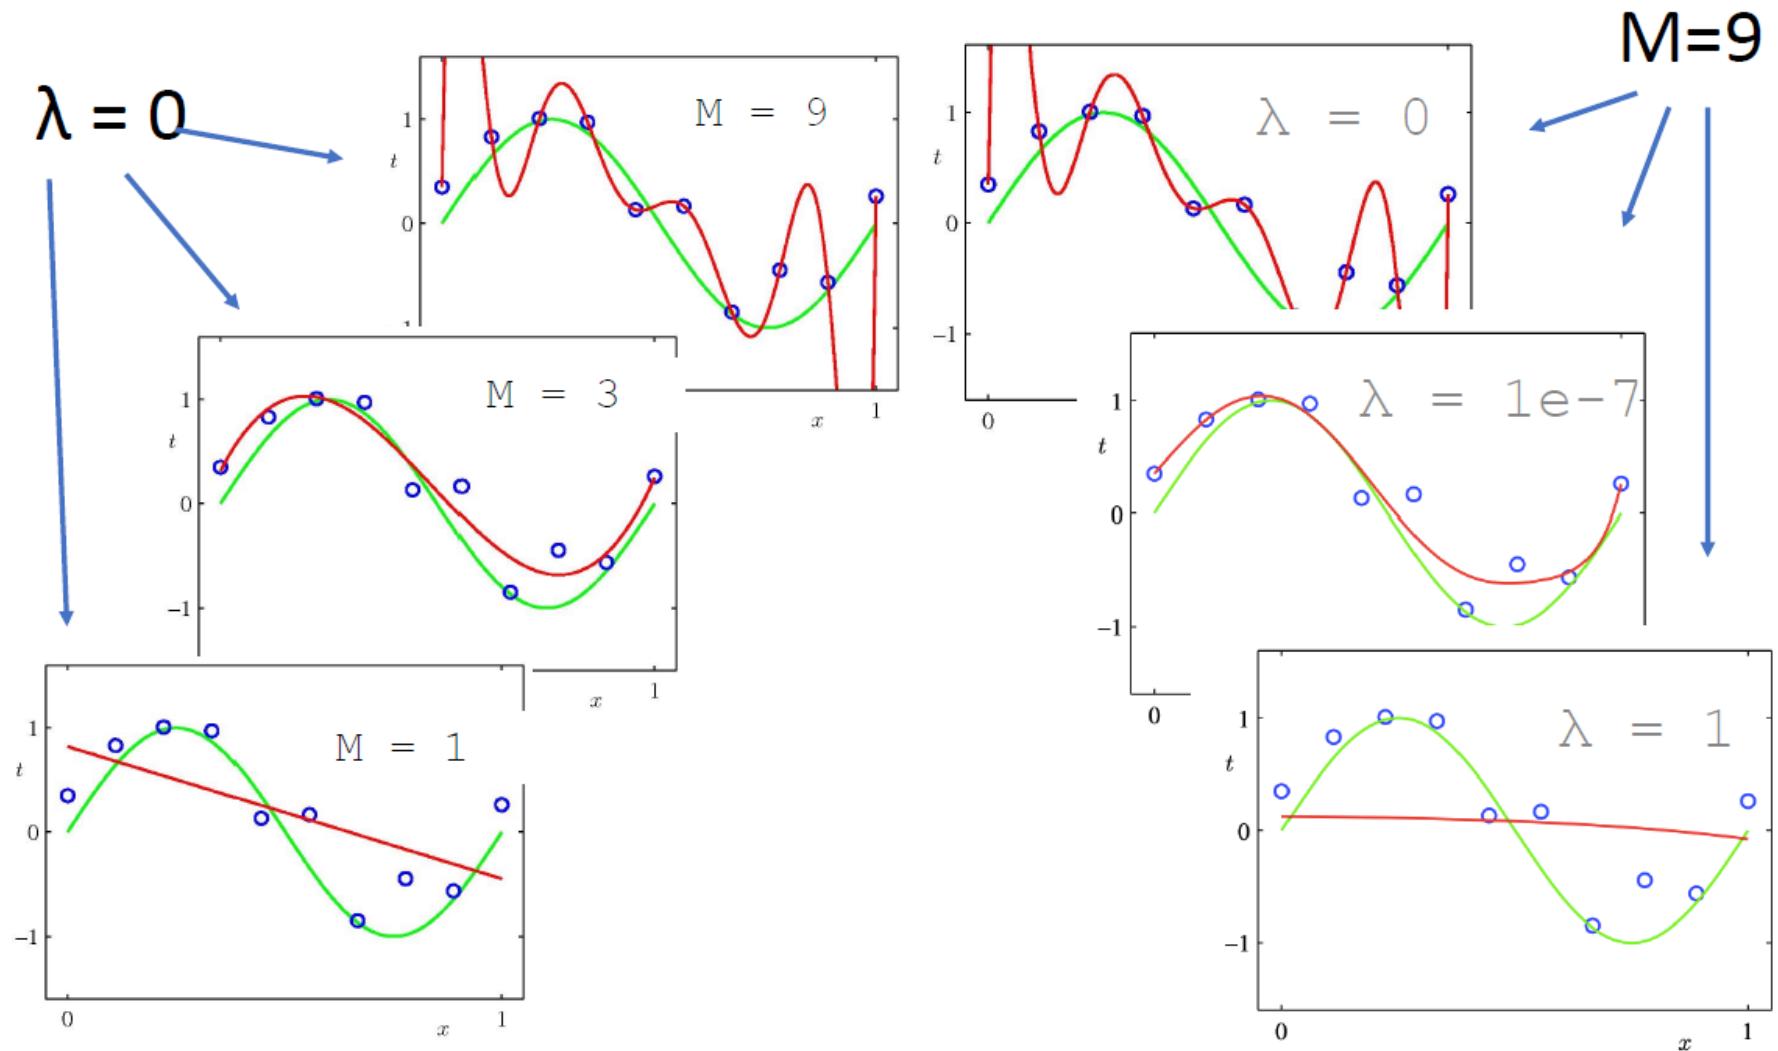
$$L(y, f(x; w)) = (y - w \cdot x)^2 + \lambda w^2$$



$$\mathbf{w} = (\lambda \mathbf{I} + \mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y}$$

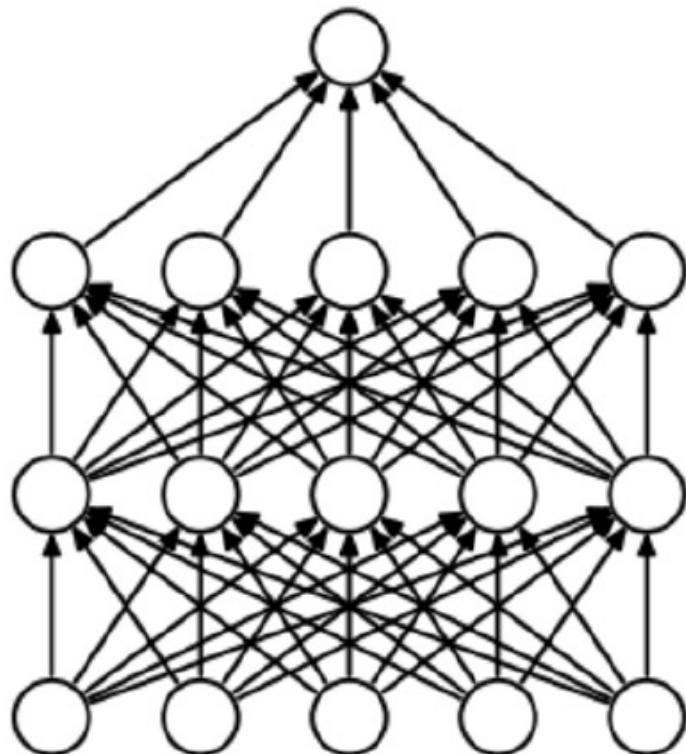


Summary of Regularization

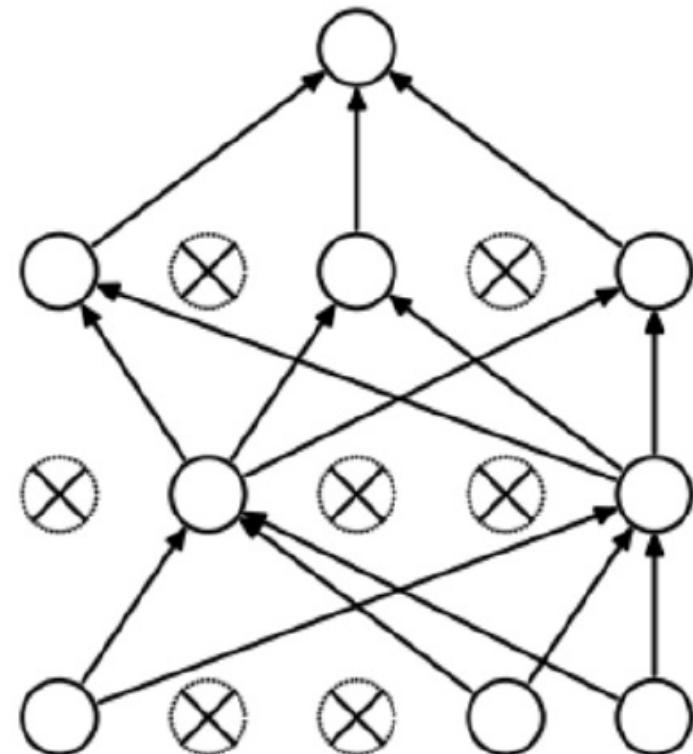


Dropout

- Instead of modifying objective function (regularization), modify the **network structure**



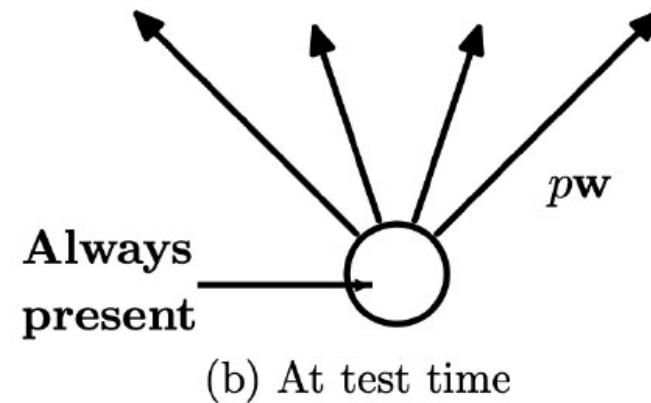
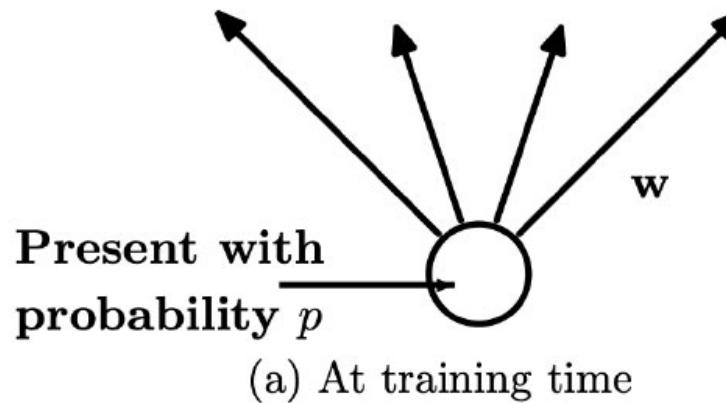
(a) Standard Neural Net



(b) After applying dropout.

Dropout

- Randomly “dropout” some neurons with probability p
- The chosen neurons change on each mini-batch
- All other learning remains unchanged



Srivastava, Nitish, et al. "Dropout: a simple way to prevent neural networks from overfitting." *Journal of machine learning research* 15.1 (2014)

Why Does Dropout Work?

- **Noise robustness**

- Dropout is similar to making random perturbations in the input
- These small changes shouldn't change the prediction for an input
- Think “reduces variance”

Why Does Dropout Work?

- Prevents neuron co-adaptation
 - Neurons, especially near the top of the network, suffer from co-adaptation
 - The neurons take on values that rely on other neurons for correct predictions
 - Multiple points of failure: any single neuron failing can invalidate the others
 - Dropout prevents co-adaptation by randomly removing some of these neurons
 - Network is forced to learn more robust features that are useful with many different random subsets of other neurons

Why Does Dropout Work?

- Ensemble training
 - Similar to random forests or boosted models: we can combine many classifiers to produce better output
 - During training, each random permutation is one of 2^n possible neural networks
 - All networks share weights so number of parameters is still $O(n^2)$
 - Dropout training is like training a collection of 2^n “thinned” networks with extensive weight sharing
 - Test time: approximate the average prediction of all networks by using a single network without dropout

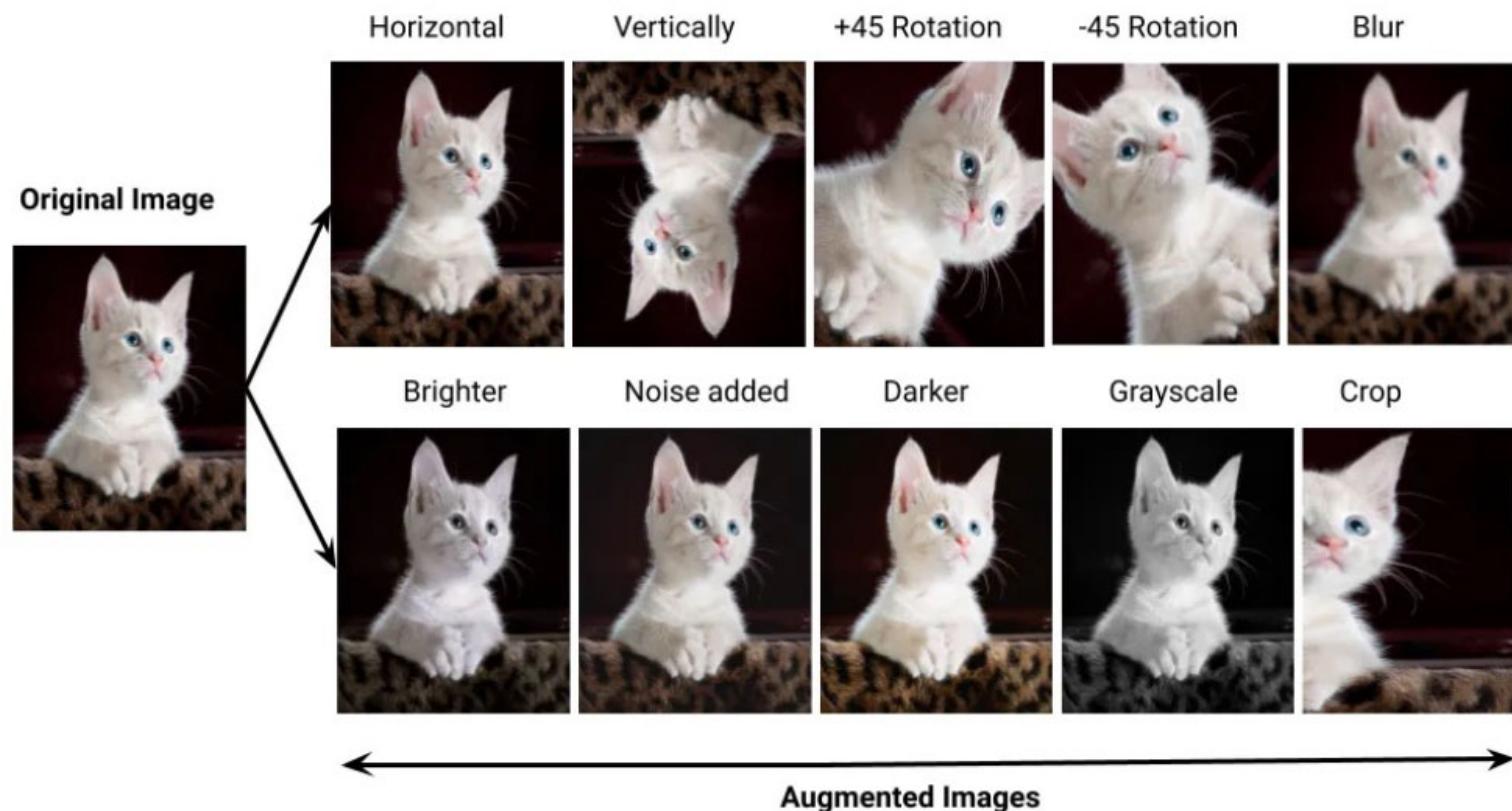
Dropout Takeaways

- Works really well for a wide variety of network types and domains
- Very easy to implement with no modifications in training
 - Some minor details (omitted) are important
- Drawbacks
 - Need to select dropout rate
 - This can depend on types of layers and network structure

Data Augmentation

- Make perturbed copies of your data that vary in ways that should not change the value nature of the output function.
- This can help prevent spurious correlations between data and output.
- Example: Distinguishing clarinet sounds from flute sounds
 - Vary the pitch of each note by + or – 1%, 2%, 3%, 4%....
 - Add background noise of different kinds and at different dB
 - Time-stretch each note a bit
 - Delay or advance the onset of the note
- This can turn 1000 data points into 100,000.

Image Augmentation



More Image Augmentation



Image Augmentation -- What if....

Your training data
looks like this?

High dynamic range
Very bright



Your testing data
looks like this?

Low dynamic range
Very dark

Image Augmentation -- What if...

- Standardize your data:
 - Make sure that you have unit variance in your batch/dataset.
- Give your data the same range overall (e.g. center your values around the same center point)
- Decorrelate your variables (can be harder for images, if every pixel is a variable)

Batch Normalization

- Normalize the data scale input to each node
- Subtract the mean value of the data
- Do this on a dimension-by-dimension basis
- Do this at every training step in gradient descent

Batch Normalization

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_1 \dots m\}$;
Parameters to be learned: γ, β

Output: $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{scale and shift}$$

Batch Normalization

Network with BatchNorm



Faster Convergence

Robust to Hyperparameters